
Data Analytics

Yong Zheng

Illinois Institute of Technology
Chicago, IL, 60616, USA



School of Applied Technology
ILLINOIS INSTITUTE OF TECHNOLOGY

Schedule

- Classification
- Classification by KNN
- Data Preprocessing



Schedule

- Classification
- Classification by KNN
- Data Preprocessing



Predictive Models We learnt

- Multiple Linear Prediction
 - One dependent variable, $y \rightarrow$ Numerical variable
 - Multiple independent variables, $x \rightarrow$ No limitations
- Linear Time-Series Model
 - One dependent variable, $y \rightarrow$ Numerical variable
 - Data with timestamp
 - AR, MA, ARMA, ARIMA
- Classification Model
 - One dependent variable, $y \rightarrow$ Nominal variable
 - Multiple independent variables, $x \rightarrow$ No limitations

Supervised v.s. Unsupervised Learning

Machine Learning Algorithms *(sample)*

	<u>Unsupervised</u>	<u>Supervised</u>
<u>Continuous</u>	<ul style="list-style-type: none">• Clustering & Dimensionality Reduction<ul style="list-style-type: none">○ SVD○ PCA○ K-means	<ul style="list-style-type: none">• Regression<ul style="list-style-type: none">○ Linear○ Polynomial• Decision Trees• Random Forests
<u>Categorical</u>	<ul style="list-style-type: none">• Association Analysis<ul style="list-style-type: none">○ Apriori○ FP-Growth• Hidden Markov Model	<ul style="list-style-type: none">• Classification<ul style="list-style-type: none">○ KNN○ Trees○ Logistic Regression○ Naive-Bayes○ SVM

Supervised Learning: Classification

- **Classification:** a supervised way to group objects
 - We must have predefined labels
 - We must have knowledge: we know some instances are labeled by predefined classes/labels/categories
- **For a Purpose of Prediction**
 - To forecast or deduce the label/class based on values of features
 - Let the machines/computers think as humans
- There are many **real-world applications**
 - Financial Decision Making, e.g., credit card application
 - Image Processing, e.g., face recognition in cameras
 - Computer/Network Security, e.g., virus or attack detection
 - Information Retrieval, e.g., relevance of a document to a query
 - Recommender Systems, e.g., rating prediction for Amazon

Classification App: Credit Card Application

First name M.I. Last name Suffix

Mailing address 1 Mailing address 2 Unit/apt.

City State ZIP code

Select the types of accounts you own. ☐ Checking ☐ Savings

Type of residence
Select One

Gross annual income
\$.00

Source of income
Select One

Employer

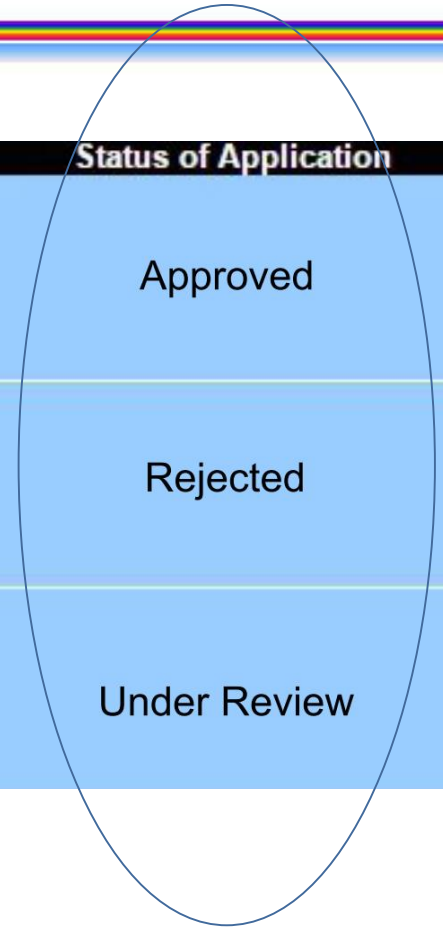
Does your credit report show any bankruptcies or seriously delinquent accounts? ☐ Yes ☐ No

Identity

Financial situation

Classification App: Credit Card Application

Date Received	Card	Status of Application
05/21/15	THE AMERICAN EXPRESS BUSINESS PLATINUM CARD	Approved
07/22/15	THE GOLD DELTA SKYMILES BUSINESS CREDIT CARD	Rejected
08/19/15	PREMIER REWARDS GOLD CARD FROM AMERICAN EXPRESS	Under Review



Classification App: Credit Card Application

Terminologies in Classification

Features					classes
Age	Gender	Status	Income	Rent	Classes
27	Female	Student	\$15,000	\$800	Approved
32	Male	Part-time	\$8,000	\$400	Rejected
29	Male	Full-time	\$50,000	\$1200	?

Knowledge (rows 1-2)

Unseen data (row 3)

Each row with features values is named as **example** or **instance**

Classification → Learn from the knowledge (examples with known labels)
build predictive models to predict the unknown examples

Classification

- Classification Tasks
- Standard Classification Process
- Evaluation: How could we know it is good or bad
- Algorithms: How to perform classification tasks



Classification

- Classification Tasks
- Standard Classification Process
- Evaluation: How could we know it is good or bad
- Algorithms: How to perform classification tasks



Classification Task

There are usually three types of classification:

1). Binary Classification

Question: Is this an apple? Yes or No.

2). Multi-class Classification

Question: Is this an apple, banana or orange?

3). Multi-label Classification

Use appropriate words to describe it:

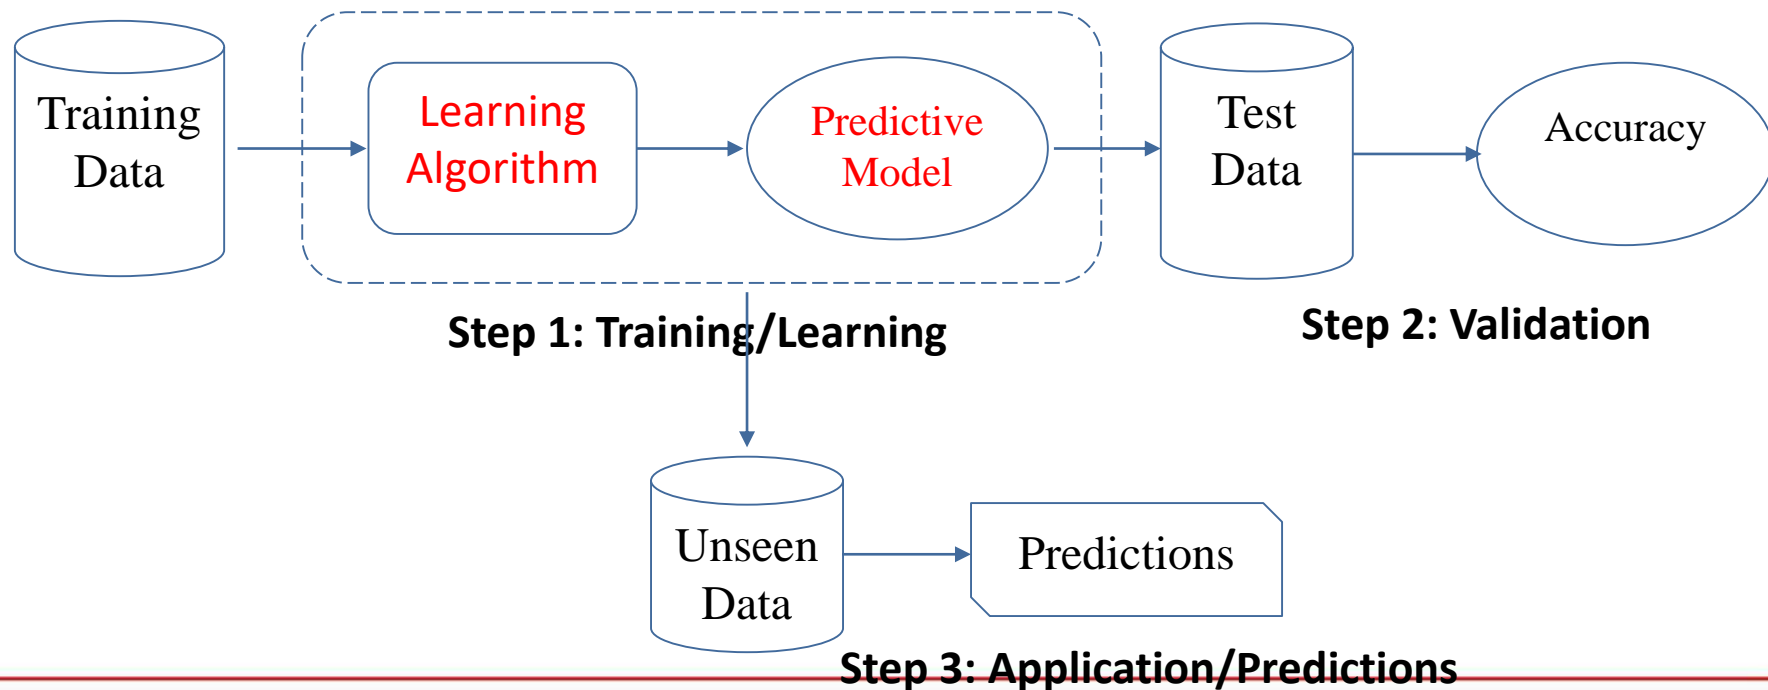
Red, Apple, Fruit, Tech, Mac, iPhone



Standard Process In Supervised Learning

- **Train:** Learn a model using the **training data**
- **Validation/Test:** Test using **test data** to assess accuracy
- **Application:** Apply the selected model to **unseen data**

$$Accuracy = \frac{\text{Number of correct classifications}}{\text{Total number of test cases}}$$



Step 3: Application/Predictions

Classification Algorithms

- Classification algorithm is the key component in the process
- They are able to learn from training and build models

There are many (supervised) classification algorithms:

- K-nearest neighbor classifier
- Naïve Bayes classifier
- Decision tress
- Logistic regression
- Support Vector Machines
- Ensemble classifiers (e.g., random forest)
- ...

Schedule

- Classification
- Classification by KNN
- Data Preprocessing



K-Nearest Neighbor (KNN) Classifier

- Problem: Identify which animal the given object it is
- Features: weights, age, gender, stripes, size, etc



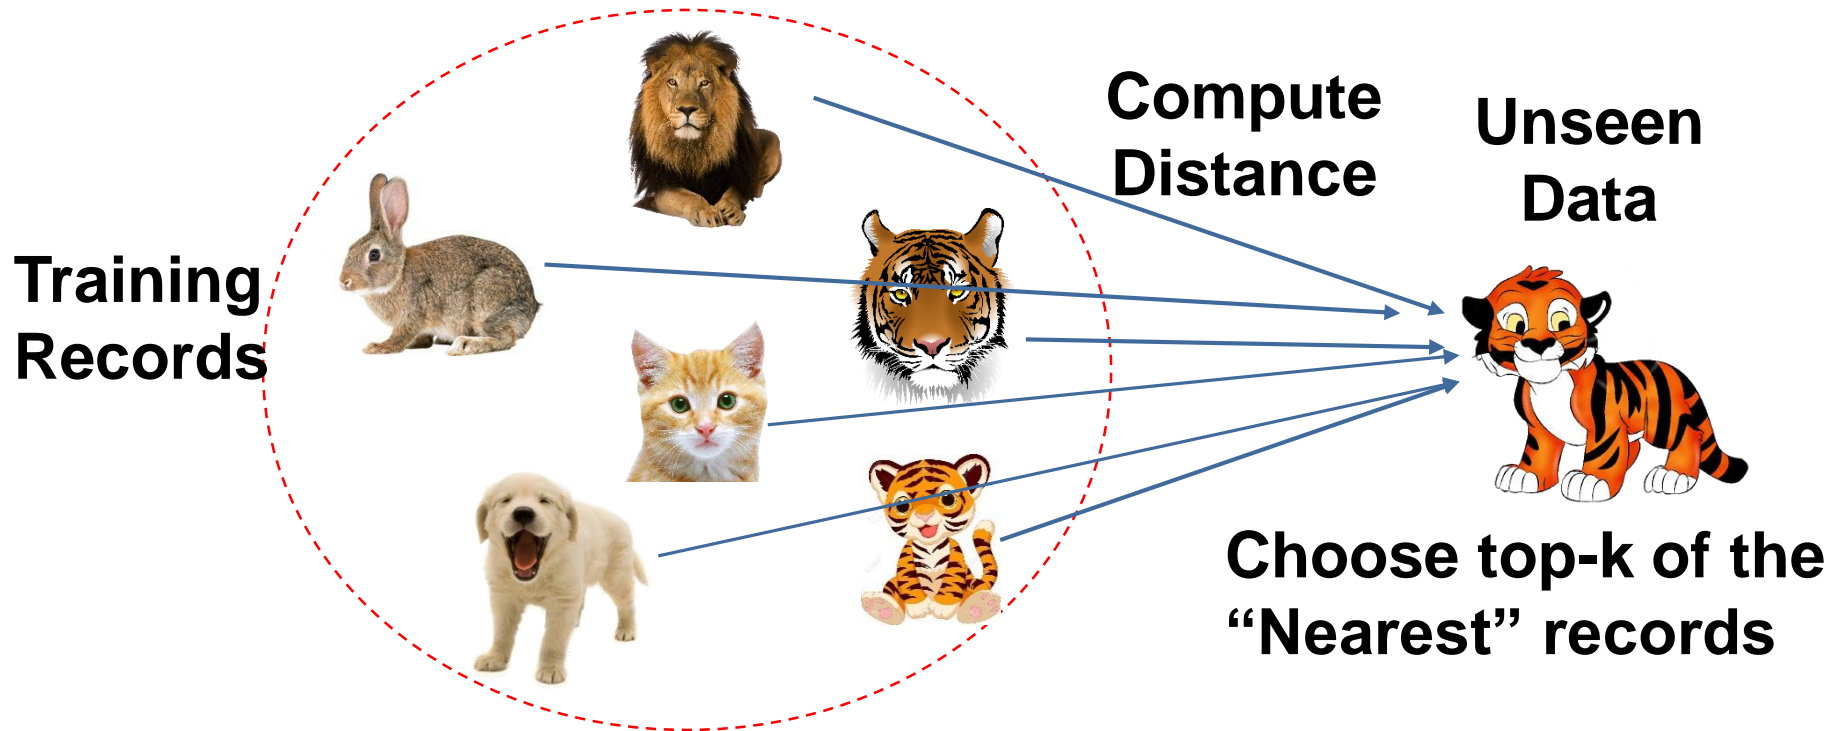
**Training
Records**

**Unseen
Data**



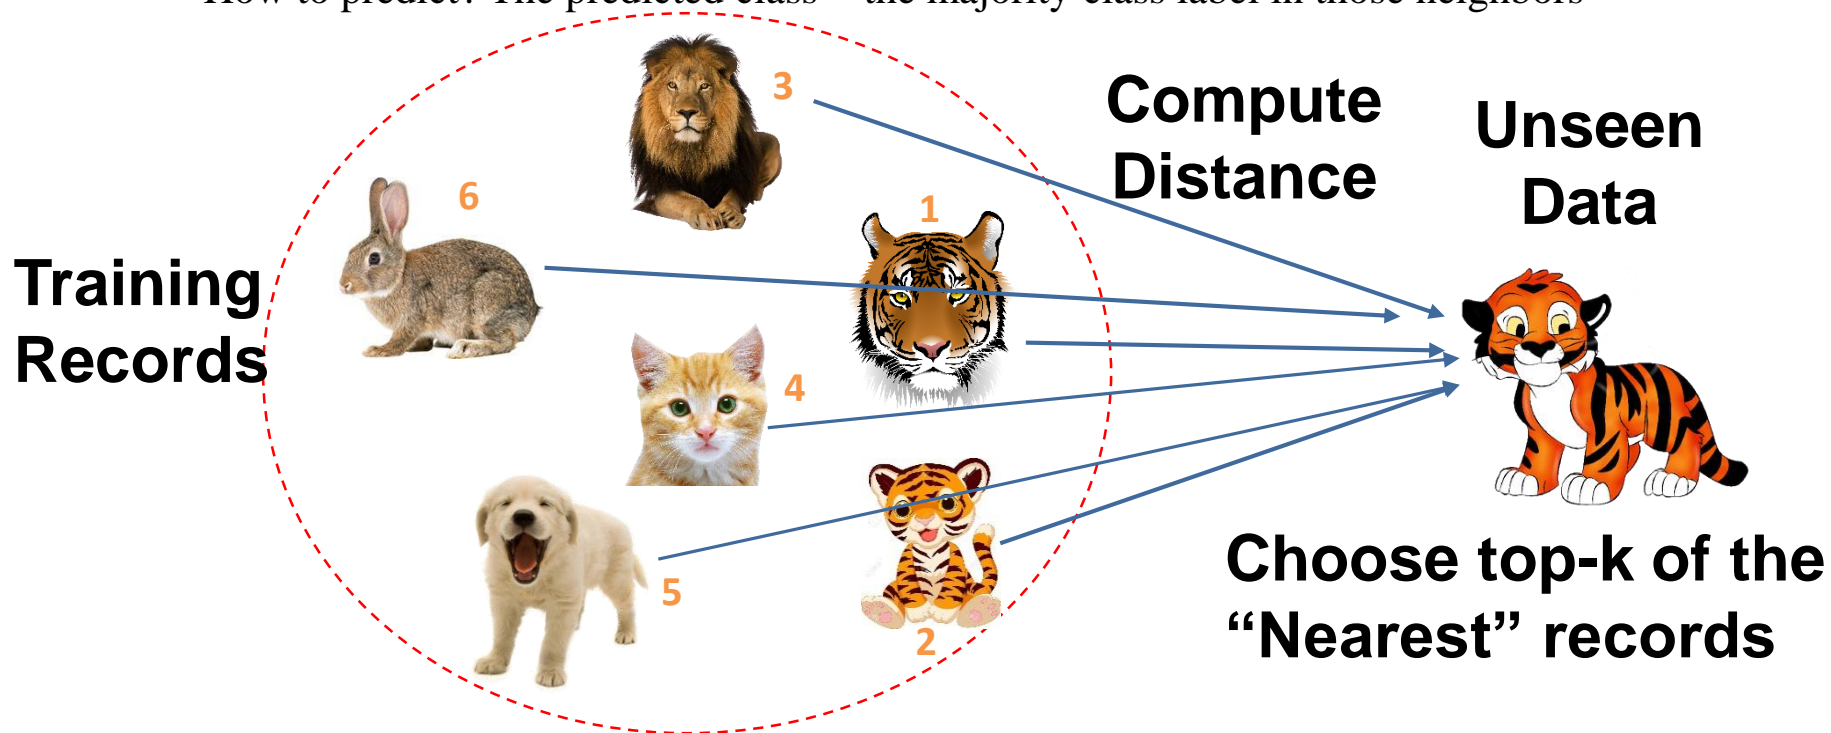
K-Nearest Neighbor (KNN) Classifier

- KNN classifier is a simple classification algorithm
- The idea behind is to classify new examples based on their similarity to or distance from examples we have seen before (in training set).



Build a KNN Classifier

- 1. Calculate distances between target and instances in train set
- 2. Identify the top-K nearest neighbor (choose an odd number for K!)
- 3. Predict labels and validate with truth
 - How to predict? The predicted class = the majority class label in those neighbors



For example, among top 3 picks ($K = 3$), 2/3 are tigers!!

Distance Measures

Assume there are n features, and two examples: X and Y .

- Consider two vectors

- ▶ Rows in the data matrix

$$X = \langle x_1, x_2, \dots, x_n \rangle$$

$$Y = \langle y_1, y_2, \dots, y_n \rangle$$

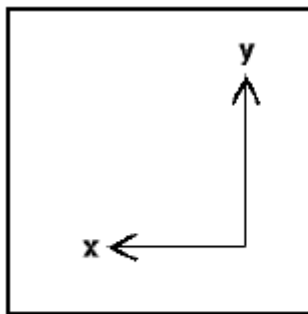
- Common Distance Measures:

- ▶ Manhattan distance: (aggregation of two right-angle legs)

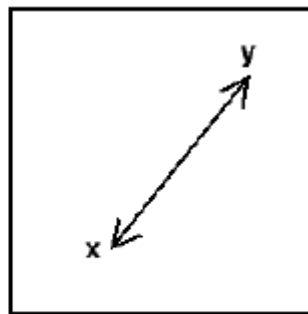
$$\text{dist}(X, Y) = |x_1 - y_1| + |x_2 - y_2| + \dots + |x_n - y_n|$$

- ▶ Euclidean distance: (length of hypotenuse)

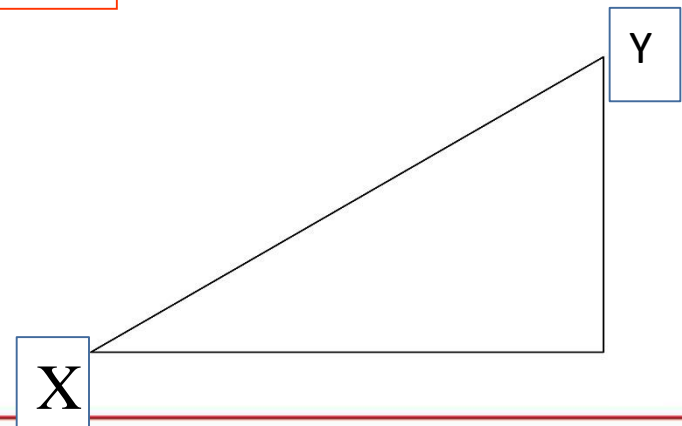
$$\text{dist}(X, Y) = \sqrt{(x_1 - y_1)^2 + \dots + (x_n - y_n)^2}$$



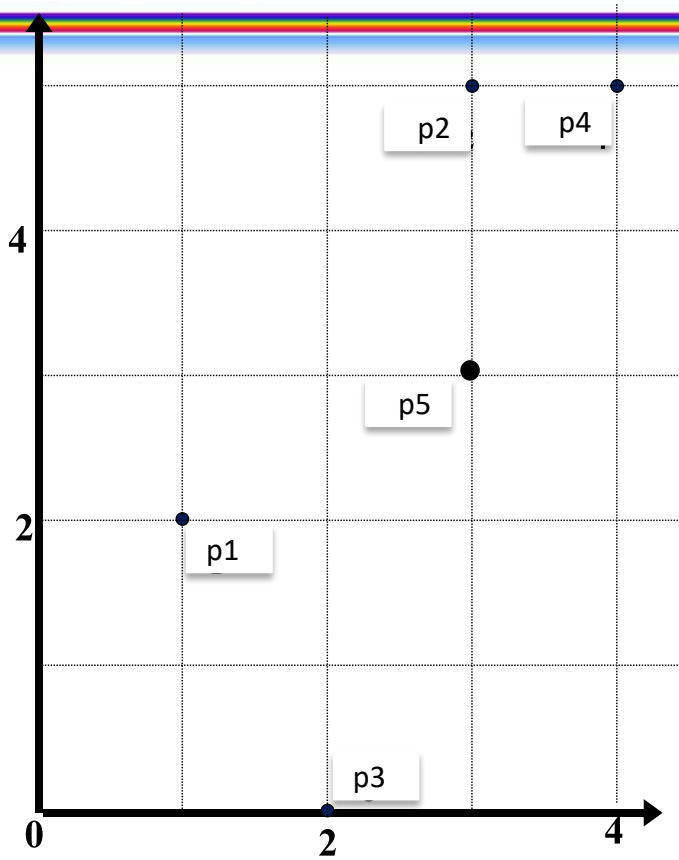
Manhattan



Euclidean



Example: Distance Measures



Data Matrix

point	feature1	feature 2	class
<i>p1</i>	1	2	Y
<i>p2</i>	3	5	N
<i>p3</i>	2	0	Y
<i>p4</i>	4	5	N
<i>p5</i>	3	3	?

Distance Matrix (Euclidean)

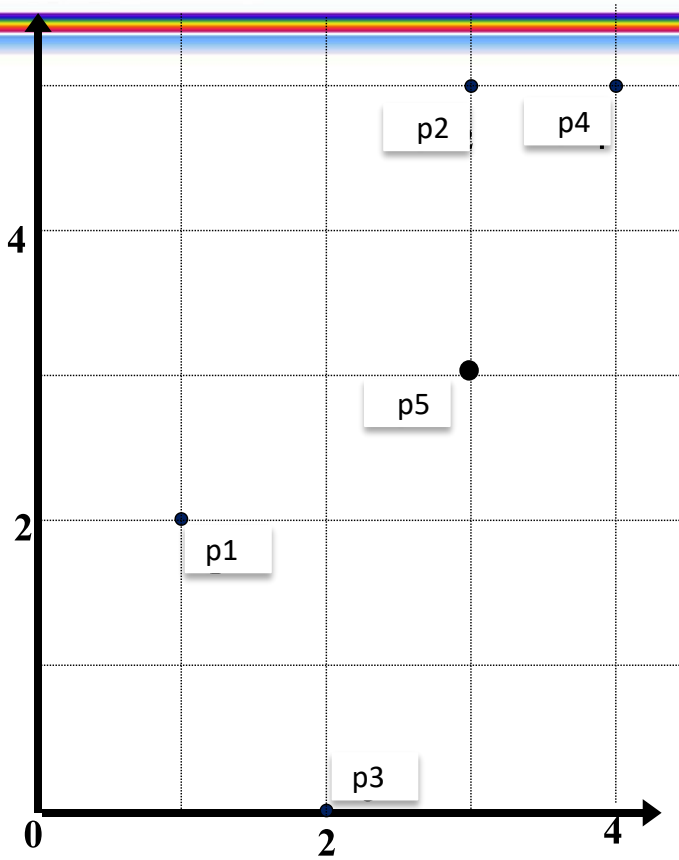
	<i>p1</i>	<i>p2</i>	<i>p3</i>	<i>p4</i>	<i>p5</i>
<i>p1</i>	0				
<i>p2</i>	3.61	0			
<i>p3</i>	2.24	5.1	0		
<i>p4</i>	4.24	1	5.39	0	
<i>p5</i>	2.24	2	3.16	2.24	0

Set K = 3

Predict label for p5

$$\text{dist}(X,Y) = \sqrt{(x_1 - y_1)^2 + \dots + (x_n - y_n)^2}$$

Time for Practice!



Data Matrix

point	feature1	feature 2	class
<i>p1</i>	1	2	Y
<i>p2</i>	3	5	N
<i>p3</i>	2	0	Y
<i>p4</i>	4	5	N
<i>p5</i>	3	3	?

Distance Matrix (Manhattan)

	<i>p1</i>	<i>p2</i>	<i>p3</i>	<i>p4</i>	<i>p5</i>
<i>p1</i>	0				
<i>p2</i>		0			
<i>p3</i>			0		
<i>p4</i>				0	
<i>p5</i>					0

Set K = 3

Predict label for p5

$$\text{dist}(X,Y) = |x_1 - y_1| + |x_2 - y_2| + \dots + |x_n - y_n|$$

Classification Algorithm: K-Nearest Neighbor Classifier

More Questions



Q1: How about categorical features?

point	feature1	feature2	class
<i>x1</i>	1	2	Y
<i>x2</i>	3	5	N
<i>x3</i>	2	0	Y
<i>x4</i>	4	5	N
<i>x5</i>	3	3	N

Color	Weight (lbs)	Stripes	Tiger?
Orange	300	no	no
White	50	yes	no
Green	490	yes	yes
White	510	yes	yes
Orange	490	no	no

Answer: Convert a categorical feature to binary features

Color	Weight (lbs)	Stripes
Orange	300	no
White	50	yes
Green	490	yes
White	510	yes
Orange	490	no



Orange	White	Green	Weight (lbs)	Stripes
1	0	0	300	0
0	1	0	50	1
0	0	1	490	1
0	1	0	510	1
1	0	0	490	0

Q2: Is feature normalization required?

Feature normalization is used to convert values in a feature to the same or similar scales with values in other features.

Answer: Yes, normalization is required, otherwise, the distance calculation will be influenced by the larger features!!!!

Orange	White	Green	Weight (lbs)	Stripes
1	0	0	300	0
0	1	0	50	1
0	0	1	490	1
0	1	0	510	1
1	0	0	490	0

Min-max normalization: transformation from OldValue to NewValue

$$NewValue = NewMin + \frac{OldValue - OldMin}{OldMax - OldMin} \times (NewMax - NewMin)$$

Summary

❑ K-Nearest Neighbor (KNN) Classifier

A simple classifier, a lazy learner

- 1). Choose an odd number for K
- 2). Calculate distances between target and instances in training set
- 3). Pick the top KNN and assign the majority label as prediction

❑ Extended Problems in Classification Algorithms

- Q1. Is it able to take categorical features? If Yes, how to treat them
- Q2. Is normalization required?
- Q3. How to alleviate overfitting problem?

Note: they are general concerns in classification, not only KNN.

R Practice

❑ German Credit Data

Download if from UCI ML Repository

<http://archive.ics.uci.edu/ml/datasets/Statlog+%28German+Credit+Data%29>

❑ Features and label

You need to read the page to understand the features and labels

Label = 1 (Good) or 2 (Bad) Credit

Features:

Existing checking account, credit history, purpose, credit amount, etc

You need to figure out feature types one by one

R Practice

❑ Load the data and take a look at it

```
setwd("C:/Users/cool/Desktop")
gc <- read.csv("germancredit.csv")
head(gc)
```

```
## Default checkingstatus1 duration history purpose amount savings employ
## 1      0           A11         6    A34    A43   1169    A65    A75
## 2      1           A12        48    A32    A43   5951    A61    A73
## 3      0           A14        12    A34    A46   2096    A61    A74
## 4      0           A11        42    A32    A42   7882    A61    A74
## 5      1           A11        24    A33    A40   4870    A61    A73
## 6      0           A14        36    A32    A46   9055    A65    A73
## installment status others residence property age otherplans housing
## 1           4    A93   A101         4    A121   67    A143   A152
## 2           2    A92   A101         2    A121   22    A143   A152
## 3           2    A93   A101         3    A121   49    A143   A152
## 4           2    A93   A103         4    A122   45    A143   A153
## 5           3    A93   A101         4    A124   53    A143   A153
## 6           2    A93   A101         4    A124   35    A143   A153
## cards  job liable tele foreign
## 1      2 A173     1 A192   A201
## 2      1 A173     1 A191   A201
## 3      1 A172     2 A191   A201
## 4      1 A173     2 A191   A201
## 5      2 A173     2 A191   A201
## 6      1 A172     2 A192   A201
```

R Practice

❑ Data Preprocessing if necessary

Convert the dependent var to factor. Normalize the numeric variables

`gc$Default <- factor(gc$Default)` → convert label to nominal data

`num.vars <- sapply(gc, is.numeric)` → extract numerical variables

`gc[num.vars] <- lapply(gc[num.vars], scale)` → normalize selected data, the function could be `normalize` or `scale`

Selecting only 3 numeric variables for this demonstration, just to keep things simple

`myvars <- c("Duration", "Amount", "Instalment")`

`gc.subset <- gc[myvars]`

R Practice

❑ Evaluate model by hold-out evaluations

predict on a test set of 100 observations. Rest to be used as train set.

```
set.seed(123)
```

```
test <- 1:100
```

```
train.gc <- gc.subset[-test,]
```

```
test.gc <- gc.subset[test,]
```

```
train.def <- gc$Default[-test]
```

```
test.def <- gc$Default[test]
```

R Practice

- ❑ User KNN to build the models and make predictions

```
library(class)
```

```
knn.1 <- knn(train.gc, test.gc, train.def, k=1)
```

```
knn.5 <- knn(train.gc, test.gc, train.def, k=5)
```

```
knn.15 <- knn(train.gc, test.gc, train.def, k=15)
```

R Practice

❑ Evaluation Based on the Accuracy

```
install.packages('Metrics', dependencies = TRUE)
library(Metrics)
accuracy(actual data, predictions)
```

```
accuracy(test.def, knn.1)
accuracy(test.def, knn.5)
accuracy(test.def, knn.15)
```

```
> accuracy(test.def, knn.15)
[1] 0.9
> accuracy(test.def, knn.1)
[1] 0.72
> accuracy(test.def, knn.5)
[1] 0.78
> accuracy(test.def, knn.15)
[1] 0.9
```

R Practice

❑ How about N-fold cross-validation?

```
install.packages('caret', dependencies = TRUE)
```

```
library(caret)
```

```
x=gc.subset
```

```
y=gc$Default
```

```
model =
```

```
train(x,y,'knn',trControl=trainControl(method='cv',number=10),tuneGrid  
= expand.grid(k = 1:10))
```

10-folds cross validation



Use K = 1, 2, 3, ..., 10



The train function is very powerful. You can use several classification methods and evaluation methods. For more details

<https://machinelearningmastery.com/how-to-estimate-model-accuracy-in-r-using-the-caret-package/>

R Practice

□ How about N-fold cross-validation?

```
> model = train(x,y,'knn',trControl=trainControl(method='cv',number=10),tuneGrid = expand.grid(k = 1:10))
> print(model)
k-Nearest Neighbors

1000 samples
  3 predictor
  2 classes: '0', '1'

No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 900, 900, 900, 900, 900, 900, ...
Resampling results across tuning parameters:
```

k	Accuracy	Kappa
1	0.594	0.04452968
2	0.596	0.04595073
3	0.638	0.07445211
4	0.634	0.05021883
5	0.648	0.04412389
6	0.659	0.08326373
7	0.659	0.04442904
8	0.672	0.06202044
9	0.681	0.06849165
10	0.671	0.04058772

Schedule

- Classification
- Classification by KNN
- Data Preprocessing



Data Preprocessing by Using R

- Replace missing values by R
- Normalization by R
- Transformation: Numerical to Categorical
- Transformation: Categorical to Numerical



Data Preprocessing by Using R

- Replace missing values by R
- Normalization by R
- Transformation: Numerical to Categorical
- Transformation: Categorical to Numerical



Data Preprocessing by Using R

- Replace missing values by R

ave

Group Averages Over Level Combinations Of Factors

Subsets of `x[]` are averaged, where each subset consist of those observations with the same factor levels.

Keywords [univar](#)

Usage

```
ave(x, ..., FUN = mean)
```

Arguments

- x** A numeric.
- ...** Grouping variables, typically factors, all of the same `length` as `x`.
- FUN** Function to apply for each factor level combination.

<https://www.rdocumentation.org/packages/stats/versions/3.5.2/topics/ave>



Data Preprocessing by Using R

- Replace missing values by R

ifelse

Conditional Element Selection

`ifelse` returns a value with the same shape as `test` which is filled with elements selected from either `yes` or `no` `TRUE` or `FALSE` .

Keywords [programming](#), [logic](#)

Usage

```
ifelse(test, yes, no)
```

Arguments

test an object which can be coerced to logical mode.

yes return values for true elements of `test` .

no return values for false elements of `test` .

<https://www.rdocumentation.org/packages/base/versions/3.5.2/topics/ifelse>



Data Preprocessing by Using R

- Replace missing values by R

##	Country	Age	Salary	Purchased
## 1	France	44	72000	No
## 2	Spain	27	48000	Yes
## 3	Germany	30	54000	No
## 4	Spain	38	61000	No
## 5	Germany	40	NA	Yes
## 6	France	35	58000	Yes
## 7	Spain	NA	52000	No
## 8	France	48	79000	Yes
## 9	Germany	50	83000	No
## 10	France	37	67000	Yes

```
dataset$Age <- ifelse(is.na(dataset$Age),  
                      ave(dataset$Age, FUN = function(x)  
                          mean(x, na.rm = TRUE)),  
                      dataset$Age)  
  
dataset$Salary <- ifelse(is.na(dataset$Salary),  
                         ave(dataset$Salary, FUN = function(x)  
                             mean(x, na.rm = TRUE)),  
                         dataset$Salary)
```


Data Preprocessing by Using R

- Replace missing values by R
- Normalization by R
- Transformation: Numerical to Categorical
- Transformation: Categorical to Numerical

Normalization by Using R

```
#extract numerical variables  
num.vars<-sapply(data, is.numeric)  
#normalize selected data using function scale  
data[num.vars] <-lapply(data[num.vars], scale)
```

```
#min-max normalization to scale [0, 1]  
data[num.vars] <-apply(data[num.vars], 2, FUN = function(x) (x - min(x))/(max(x)-min(x)))
```



Index for rows or columns
2 = apply function based on columns
1 = apply function based on rows

difference among apply(), lapply(), sapply(), tapply()?

<https://www.guru99.com/r-apply-sapply-tapply.html>



Data Preprocessing by Using R

- Replace missing values by R
- Normalization by R
- Transformation: Numerical to Categorical
- Transformation: Categorical to Numerical



Data PreProcessing

❑ Convert Numerical variable to Nominal variable in R

F1	F2	F3	F4	Class
C3	0	0	2	—
C2	1	0	5	+
C1	0	1	8	—
C2	1	1	16	—
C1	1	0	23	+
C3	0	1	11	+

Usually we use the cut function to create N groups

Data = cut(dataColumn, N)

```
> data=read.table("book1.csv", head=T, sep=',')
> data[,4]
[1] 2 5 8 16 23 11
> data[,4]=cut(data[,4], 3)
> data[,4]
[1] (1.98,9] (1.98,9] (1.98,9] (9,16] (16,23] (9,16]
Levels: (1.98,9] (9,16] (16,23]
> head(data)
  F1 F2 F3      F4 Class
1 C3  0  0 (1.98,9]    -
2 C2  1  0 (1.98,9]    +
3 C1  0  1 (1.98,9]    -
4 C2  1  1  (9,16]    -
5 C1  1  0 (16,23]    +
6 C3  0  1  (9,16]    +
```

Data Preprocessing by Using R

- Replace missing values by R
- Normalization by R
- Transformation: Numerical to Categorical
- Transformation: Categorical to Numerical



Data PreProcessing

❑ Convert Nominal Variable to Dummy variables in R

F1	F2	F3	F4	Class
C3	0	0	2	—
C2	1	0	5	+
C1	0	1	8	—
C2	1	1	16	—
C1	1	0	23	+
C3	0	1	11	+

```
install.packages("dummies")  
library(dummies)  
data=read.table("book1.csv", head=T, sep=',')  
df=dummy.data.frame(data,names=c("F1"))
```

```
> df=dummy.data.frame(data,names=c("F1"))  
>  
> df  
  F1C1 F1C2 F1C3 F2 F3 F4 Class  
1    0    0    1  0  0  2    —  
2    0    1    0  1  0  5    +  
3    1    0    0  0  1  8    —  
4    0    1    0  1  1 16    —  
5    1    0    0  1  0 23    +  
6    0    0    1  0  1 11    +
```

Note that it will create N dummy variables if there are N values in the nominal variable

Midterm Exam

- Time: Mar 26, 08:30 AM – 09:50 AM
- Location: SB 111
- Closed Note, Closed Book, Closed Devices
- **You can bring a calculator.** You can NOT share calculator with others.
- The questions in the exam will be the similar ones in your assignments; but you do not need to produce the R outputs, the outputs will be given in the exam papers.

